Machine Learning to the rescue

Jan Macháček

May 30, 2018

Abstract

Machine learning to the rescue!... the first question is "to the rescue of what?"; immediately followed by "when is it indeed rescued?". The answers to these questions are crucial; luckily, the software engineering process is quite used to asking and answering these questions. Careful project analysis and inception, followed by continous integration and continous deployment in development (supported by adequate tests); all overseen by systematic project governance leads to successful software projects. This paper's proposition is that machine learning projects that are to apply established machine learning approaches and algorithms are no different than any other software project; and they must follow all practices of software engineering.

1 Software engineering

...

The Joel test[4] is a product of one man's biased, ad-hoc, informal, ... view of what makes successful software projects. Annoyingly, high scores on the Joel test correlate with successfully delivered software projects. Spolsky acknowledges that it's possible for a small team of cowboys to deliver amazing software with score of 0, as much as it is possible for a team that scores 12 to be the software equivalent of the Titanic. Nevertheless, high scores on the test correlate with good practices and discipline, which usually leads to good software.

This paper's proposition is that "business" projects that use machine learning are no different than any other software project; and that all practices of software engineering have to be applied to the machine learning subsystems. Specifically, that the 12 points on the Joel test are just as applicable, but with additional 12 points.

1. Version control for the ML models and data sets used to train them
   Joel: Do you use source control?

2. Single-step / automated data selection, model training, evaluation, and deployment
   Joel: Can you make a build in one step?

At least daily training and deployment process
   Joel: Do you make daily builds?

3. The results of the BI queries that humans process define what ML should solve
   Joel: Do you have a bug database?

4. Versioned, testable; continuously tested and sanity-checked BI
   Joel: Do you fix bugs before writing new code?

5. BI that allows any query to be answered in under 10 minutes

Joel: Do you have a spec? Do you have an up-to-date schedule? Do programmers have quiet working conditions? Do you use the best tools money can buy? Do you have testers? Do new candidates write code during their interview? Do you do hallway usability testing?

6. Versioned, testable; continuously tested and sanity-checked analytics (BI)

7. Monitoring on the BI environment to identify queries that use normalised data

8. The results of the BI queries that humans process define what ML should solve

9. Ingestion components decoupled from the rest of the system

10. Versioned, testable; continuously tested and sanity-checked data sets

11. Pre-computed "return constant" model

12. Versioned, testable; continuously tested and sanity-checked model storage with training and validation data set references

13. Model deployer and "debugger"

1. Versioned, testable; continuously tested and sanity-checked analytics (BI)

2. Any BI query can be answered under 10 minutes

3. Monitoring on the BI environment to identify queries that use normalised data

4. The results of the BI queries that humans process define what ML should solve

5. Ingestion components decoupled from the rest of the system

6. Versioned, testable; continuously tested and sanity-checked data sets

7. Pre-computed "return constant" model

8. Versioned, testable; continuously tested and sanity-checked model storage with training and validation data set references

9. Model deployer and "debugger"

But where is the ML that builds the model? That's the code that the engineering teams need to build to replace item 6.

Once the first four steps are known, the engineering teams can implement the remaining steps of the pipeline. If the system that is to take advantage of ML is event-based, the event delivery mechanism provides the decoupling, resulting in architecture shown in Figure 1.


Figure 1: ML pipeline in event-based system

If the front-end system is not event-sourced, the ingestion must be decoupled using a read-only replica of the live data. Notice in Figure 2 the flow of the data: the data is pushed into the read-only replica in the first step to allow the front-end system to control the load on its data store; from the read-only replica, the data is pulled into the ML data store.

Regardless of the approach used (or even if a hybrid approach is deployed), the entire system has to be aware of any back-pressure.

The ML team maintains the tooling for the pipeline, consults on the best models, researches, ...; but the product teams (that ultimately work on the service that *uses* the model) have the first dibs on implementing the model. Successful implementation of this strategy means that anyone can implement a new model (even if only to just see what will happen!), train it, debug it, and deploy it all within a single day. All the mechanics of data ingestion, storage, versioning; runtime of training a model, evaluation, storage, versioning; debugging and deploying; and the usage is all implemented.


Figure 2: ML pipeline in non event-based system

In this sense, the machine learning code is just like any other ordinary code; it is subject to all the high engineering standards and safeguards.

2 Implementation choices

Attach technologies to the blocks in Figure 1 and Figure 2. Show on practical example, attempt to pull out reusable blocks and show examples of good and bad code. Think the recent ML experiments; demonstrating just how much time the initial research and experiments take. Reinventing the experimental wheel must be avoided–the ML team should curate the bootstrapping environment, making it available "on demand" for other teams. The work on the ML core code can only successfully happen when all other pieces are in place.

2.1 Initial research

The practical application of the "ML enablement process" aimed to deliver failure predictions and automatic error recovery based in an event-based architecture system. The system already published seemingly too many events, but the detail in the events were very useful in building the data sets for the ML project. The first step the team needed to take was to build traditional business intelligence database– the goal was to find out whether our human knowledge and experience allowed us to find meaningful information in the data.

The system's journal (Apache Kafka[1]) is configured with 7 days' message retention policy; at the very start of the work, the team was able to download events for the last 7 days. This was approximately 60 GiB of uncompressed Protocol Buffers[3] binary messages. Because of the compactness of the binary wire representation of the Protocol Buffers messages, this translated to 60 GiB of a RDBMS (MySQL[2]) storage requirements. The loader program performed batch JDBC inserts into the MySQL tables. The

HttpRequest

correlationId	varchar(36)
time	timestamp
method	varchar(10)
uri	varchar(255)
headers	text
entityContentType	varchar(64)
entity	blob

HttpResponse

correlationId	varchar(36)
duration	integer
statusCode	integer
headers	text
entityContentType	varchar(64)
entity	blob

Status

correlationId	varchar(36)
hostId	varchar(50)
serviceId	varchar(50)
time	timestamp
state1	integer
state2	integer
lastState1	integer
lastState2	integer
entity	blob
lastEntity	blob

Figure 3: Database schema

events in the $HttpRequest$ and $HttpResponse$ topics were imported to the matching tables, the events in the $ServiceStateChanged$, $ServiceStateUnchanged$, $ServiceAdded$, and $ServiceRemoved$ were pre-processed and inserted into the $Service$ table. The tables are not even in the 1st normal form; there is only the primary key, but there are no foreign key constraints; during the batch import, there are no indexes.

The first seven days' worth of data import ended up as 5 103 397 rows in $HttpRequest$, 5 103 397 rows in $HttpResponse$ [1], and 54 566 575 rows in $Service$. This imported database was the starting point for the human analysis code. "How many $statusCode \neq 200$ in $HttpResponse$ are there?"–(72 796); "what is the average response duration?"–(4.0759 ms); "what is the average response duration for successful responses"–(4.9862 ms), and so on. In order to answer these questions, the database needed a few indexes: remember, "any ad-hoc analytics query needs to be answerable within 10 minutes." The db $.m4.4xlarge$ instance was completely sufficient for the first 7 days' worth of data; even for the first 70 days' worth

---

[1] The same number of rows in the $HttpRequest$ and $HttpResponse$ is a very good sanity-check!

of data. It delivered just over $150\,\mathrm{MiB\,s^{-1}}$ in read and write throughput, and with only a few users running queries, its CPU usage never got above 50 %; it was allocated 5 TiB of SSD storage.

The first valuable answer came from the analysis of the frequency of different errors by joining $request$ and $resonse$ on $correlationId$ where $statusCode \neq 200$, grouping and counting by $response.entity$. The top 3 errors turned out to be fairly simple to fix, and the fixes resolved 90 % of all observed errors. Beyond the top 3, the errors had no trivial cause and no trivial fix; the team suspected that the errors were cause by a combination of parameters passed in the requests, the previous state of the device doing the work, the network conditions, ...; something that could not be trivially reproduced nor fixed. Interestingly, the 3 least frequent errors did turn out to have trivial fixes: in the sense that there was no fix, but better failure reporting in low or missing resource conditions, misconfiguration, etc.

These were great first results. The next valuable feature was to predict a failure before it actually happened. This would reduce the pressure & stress in the operations team.

uilding time series of requests, responses, and status changes. The tool was simple SQL $create\ table\ ...\ from\ select\ ...\ inner\ join$; the sequence of events ended up as 50 182 398 rows, describing 545 357 sequences[2]. The data had sequences such as $(req_1, res_1) \rightarrow \mathcal{I}dle, (req_2, res_2) \rightarrow \mathcal{B}usy, \ldots, (req_n, res_n) \rightarrow \mathcal{E}rror$, or even $(req_1, res_1) \rightarrow \mathcal{I}dle, (req_2, res_2) \rightarrow \mathcal{B}usy$ (without further requests) $\mathcal{E}rror$. This work implements the first part of the overall diagram (viz Figure 4), and was able to satisfy (TODO) X, Y, and Z items on the ML readiness test.

Nevertheless, this initial analysis was sufficient to build data sets for failure prediction. the team deployed code more sophisticated than $select\ from\ ...\ group\ by\ ...\ having\ ....$

2.2 Buy vs. build (AWS services vs. custom code)

2.3 Testing

3 Compromises

How to manage the growing data sizes and time to train the models; particularly with complex models that need a lot of computation to train? Where to keep the old data and old models? Is there ever value in digging out models and data that are 10 versions old?

---

[2] Sanity-check: there are periodic health-checks for every service contributing nearly 80 % of the HTTP requests and responses, and a typical sequence is made up of a single request and response, and multiple (usually > 10) service changed rows; there cannot be more sequences than requests and responses. The 545 357 number of sequences is therefore a sane number of sequences, given the sizes of the tables
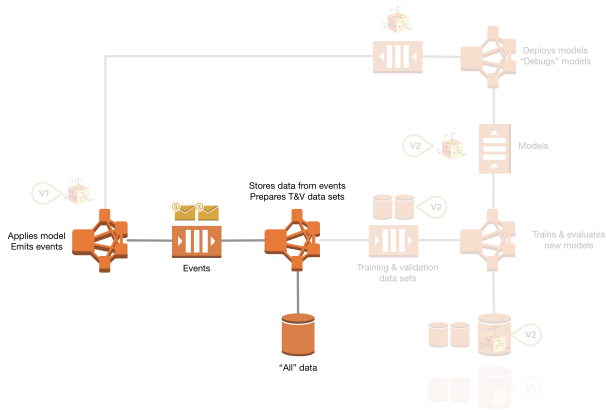
Figure 4: Data collection and BI

References

[1] Apache Kafka. *Apache Kafka*. `https : / / kafka . apache.org`.

[2] *MySQL*. `https://www.mysql.com/`.

[3] *Protocol Buffers*. `https://developers.google.com/ protocol-buffers/`.

[4] Joel Spolsky. *The Joel Test: 12 Steps to Better Code*. `https : / / www . joelonsoftware . com / 2000 / 08 / 09 / the-joel-test-12-steps-to-better-code/`.