Scala Tutorial I

Jan Macháček

June 12, 2018

Abstract

Scala is a fusion language that combines functional and object-oriented programming paradigms in a syntax that is similar to most other *C-like* languages. The ...

1 Syntax crash-course

Scala's syntax follows the syntax of other C-like languages, though–like Pascal–the type specification follows an identifier. Scala's `class` behaves exactly like Java's `class`, and its syntax is not wildly different. (See Listing 1.)

```
Class declaration means the same thing as Java; constructor parameters are specified in the block immediately following the class name.
Note that the types follow the identifier; instead of String constructorParam1 Scala uses constructorParam1: String
class MyClass(constructorParam1: String, constuctorParam2: Int) {

  Methods begin with the keyword def, followed by name and parameters. The return type follows similar pattern; Unit means void.
  The body of the method follows the equals sign.
  def execute(methodParam1: List[Int]): Unit = {

  }

}

Constructing instances uses the typical new keyword...
new MyClass("foo", 42)
  Method invocation is exactly like Java's
  .execute(List(1, 2, 3))
```

Listing 1: Classes and methods

Most of this syntax is familiar and unsurprising; the only thing that might feel odd is the square bracket for "generics" in `List[Int]`: in Java, this would be written as `List<Int>`. This is part of Scala's legacy. A long, long time ago, XML was very exciting; and Scala allows XML literals. These XML literals use the angle brackets. This meant that a different symbol had to be used for type parameters. Because the square bracket is used for type parameters, array indexing is also done using regular parentheses[1].

Interfaces use the `trait` keyword in Scala; their usage and features are similar to `interface` in Java (particularly Java 8 which adds default implementations). It is possible to make anonymous implementations of a `trait`, as well as to implement it in ordinary `class`es. (Viz Listing 2.)

```
Apart from the trait keyword, the syntax is unremarkable
trait ReportGenerator {
  Interface methods are public and abstract; they specify parameters and return type
  def generate(userId: Int): Array[Byte]
}

A trait can be implemented in a class using the keyword extends. Additional traits to be implemented use the with keyword.
class ReportGeneratorImpl extends ReportGenerator with Cloneable {
  Instead of the @Override annotation Scala uses the override keyword.
  override def generate(userId: Int): Array[Byte] = ...
  override def clone(): AnyRef = ...
}

It is also possible to make an anonymous implementation of a trait using the new keyword.
new ReportGenerator {
  override def generate(userId: Int): Array[Byte] = ...
}
```

Listing 2: Traits

[1] I know, it looks like VisualBasic or Fortran (formerly FORTRAN)!

It is worth noting that there is no special syntax for *array of X* in Scala. Instead, it uses `Array` with the specified type parameter. (So, Java's `byte[]` becomes `Array[Byte]`, `User[]` becomes `Array[User]`, and so on.) Also notice the `AnyRef` in the implementation of the `clone()` method–it is equivalent to `java.lang.Object`.

So far, there are no major surprises: classes and interfaces work just like most other languages, constructor, method, and parameter definition also looks fairly ordinary. The syntax for functions (finally!) uses parameters (each with its type following the `:` symbol, if needed) followed by fat arrow ⇒, and the body of the function. Usage is the same as Java 8; and Scala's collections library contains concepts that are fairly similar.

```
Double every Int in the list
List(1, 2, 3, ...).map { (x: Int) ⇒ 2 * x }

Scala can infer the type of x, so there's no need to specify it
List(1, 2, 3, ...).map { x ⇒ 2 * x }

An alternative syntax uses paretheses in place of brackets
List(1, 2, 3, ...).map(x ⇒ 2 * x)

It is possible to avoid having to declare the parameter x, and use _ instead. Scala compiler replaces the every occurrence of _ with a fresh parameter
    declaration.
List(1, 2, 3, ...).map(2 * _)
```

Listing 3: Fields & variables

The last form `2 * _` is somewhat unusual and worth exploring further. The Scala compiler mechanically translates the each underscore in the function body into a parameter (and "replaces" the underscore with that parameter); `2 * _` is translated into `p_0 ⇒ 2 * p_0`. Now, exploring the functions on type `Int`, the type of the function `*` is `Int ⇒ Int`. (It is defined as `def *(that: Int): Int`). The type of `2 *` is therefore `Int ⇒ Int`, just like the type of `2 * _`, or `_ * 2`, or `x ⇒ x * 2`–assuming the parameter can be inferred to be `Int`. Consequently, it is possible to leave out the underscore in the `map` function and only write `List(1, 2, 3, ...).map(2 *)`[2].

Finally, fields (and variables) use the keywords `val` and `var`. The first declares an immutable variable (and a getter if field), the second declares a mutable variable (and a getter and setter if field); see Listing 4.

```
Field definitions have to specify initial value; use _ for default value.
class User {
  var id: Long = _
  var name: String = _
  var dob: Date = _
}

val user = new User
```

Listing 4: Fields & variables

This code is terrible! The default value for reference types (`AnyRef`s) is `null`, and variants of *zero* for primitive types. The `name` and `dob` fields in the `user` instance are `null`, and `id` is `0L`. What's worse, the class has setters for these fields, and it's possible to invoke them. The syntax is somewhat nicer–it looks like plain assignment using the `=` operator, though it is actually invoking a setter–but the semantics of the code in Listing 5 is terrible.

```
Field definitions have to specify initial value; use _ for default value.
class User {
  var id: Long = _
  var name: String = _
  var dob: Date = _
}

val user = new User
user.id = 5
user.name = "Foo Quux"
user.dob = ...
```

Listing 5: Fields & variables II

While syntactically valid code, it is very confusing. The `user` variable is declared immutable, yet it is possible to invoke its setters. The equivalent Java code would declare the `user` variable `final`, but then still use the setters to mutate it.

2 Killer features

The syntax (and its application) so far looks just like Java with less typing. There must be something else that makes it worth leaving the creature comforts of Java.

[2] This might be somewhat familiar to Java 8 programmers with method references with the double colons

- everything is an expression
- case classes
- pattern matching
- for comprehensions
- implicits
- rich type system

2.1 Everything is an expression.

In Java, C, and similar languages, there are statements and expressions. Statements do not have value, they are typical control-flow constructs. For example, in Java, `if (a == b) X else Y` cannot be assigned to a variable, because `if-then-else` is a statement. In Scala, everything yields a value that can be assigned to a variable.

2.2 Case classes.

A case class is just like a class in that it is a container for data and methods, but the fields it contains only have getters. (Immutability only goes as far as immutability of the references. Even in Scala, immutability without any additional code is equivalent to using `final` in Java.) Nevertheless, case classes are fantastically convenient to define data structures. Consider the `Person` case class defined in Listing 6.

```
case class Person(firstName: String, lastName: String, age: Int)
```

Listing 6: Case class `Person`

This is all it takes to define an immutable structure (with the T&Cs from above) with the fields `firstName`, `lastName`, and `age`; but also with appropriate `toString`, `hashCode`, and `equals` implementations. These automatically generated implementations delegate to the `toString`, `hashCode`, and `equals` methods of all the fields, in the order in which they are specified.

To create an instance of a case class, do not use the `new` keyword; instead, write the parameters directly after the case class name, as shown in Listing 7.

```
Notice that there is no new keyword; the parameter values are applied directly after the case class name.
val fq = Person("Foo", "Quux", 42)

To access the fields, use the familiar . notation.
fq.firstName   | "Foo"

Invoking the toString method prints a reasonable representation of the case class.
fq.toString   | "Person(Foo,Quux,42)"

It is possible to vary the order of the parameters if the parameter names are also specified. This can help readability.
val fb = Person(lastName = "Baz", firstName = "Foo", age = 50)
fb.toString   | "Person(Foo,Baz,50)"

Equality is implemented by delegation to the parameters' hashCode and equals methods.
Person("Foo", "Quux", 42) == fq     | true, even though they are different instances.
fq == Person("Foo", "Quux", 42)     | true, even though they are different instances.
```

Listing 7: Using case class `Person`

The consciseness of Scala's syntax is beginning to show. It would have been much more cumbersome to implement all this (including correct `hashCode`, `equals`, and reasonable `toString`) in Java. Even with correct implementations, it would not have been possible to use `==` to test for instance equality.

> **Exercise: Using case classes and functions**
> Create a new case class `Person` file in `src/main/scala/com.acme` from Listing 6. Then use `List.fill(N)(Person(...))` and useful methods in `scala.util.Random` to generate N random `Person` instances. Assign those to a variable.
>
> - find the oldest person in the list. (Hint: use `map`, `max`, and `find` functions on the `List[Person]`.)
> - find the 10 youngest people in the list. (Hint: use `sortBy` or `sortWith`, followed by `take`.)
>
> ...

```
object PersonMain extends App {

  List.fill(100) ...

}
```

Listing 8: Pattern matching

Case classes are also called *product* types. The word product refers to the possible number of values that a case class can hold. Take `case class ABC(a: Boolean, b: Boolean, c: Boolean)`: there are eight possible values: `ABC(false, false, false)`, ..., `ABC(true, true, true)`. The value 8 is the result of *multiplying* the possible values of all parameters. `Boolean`s have two values; three `Boolean` values yield $2 \times 2 \times 2$ possible values of the *ABC* type. Similarly, `case class IB(i: Int, b: Boolean)` has 4294967295×2 possible values.

Imagine, for a moment, that `Boolean` is not a built-in primitive. It would be defined as the *sum* of two [degenerate] products of 1 value: `sealed trait Boolean; case object True extends Boolean; case object False extends Boolean`. The products `True` and `False` have exactly one value; consequently, the sum type `Boolean` has $1 + 1$ possible values. Languages like Haskell[3], F#[2], and others have convenient syntax for sum types; see Listing 9.

```
Defines the data type Boolean as a sum of two products, but each product only has one value.
data Boolean = True
             | False

Defines the subscription type as either a "one-off" that carries a particular end-date, or a recurring
subscription that carries the start date and the billing period.
data Subscription = OneOff Date
                  | Recurring Date Period


sealed trait Boolean
case object True extends Boolean
case object False extends Boolean

sealed trait Subscription
case class OneOff(endDate: Date) extends Subscription
case class Recurring(startDate: Date, period: Period) extends Subscription
```

Listing 9: Sum types

Finally, using the keyword `sealed` specifies that the trait cannot be extended outside the current source file. This allows the Scala compiler to verify that pattern matches (using the `match ... case` construct) cover every possible case, reporting a warning if a case is missed.

Sums of products are very useful way to express alternatives that might not have anything in common other than being of the given type, without clumsy `instanceof` checks with pattern matching.

2.3 Patten matching.

A pattern match is a way to check that a value has the right "shape", and to pull out some or all values from that shape. Think of the simplest pattern matching expression as Java's `switch` statement in Listing 10.

```
This is the equivalent of Java's switch statement, with multi-case and a default case.
Random.nextInt(10) match {
  Matches values 0..3
  case 0|1|2|3 ⇒
  Matches only value 6
  case 6 ⇒
  Matches all other values
  case _ ⇒
}
```

Listing 10: Pattern matching

It is possible to match on much more complicated structures; such as tuples. A tuple is a collection of a specific length with elements of arbitrary types. For example `(1, "foo")` is a tuple of 2 elements of types `Int` and `String`; `("foo", "bar", 1.3)` is a tuple of 3 elements of types `String`, `String`, and `Double`. It is possible to pattern match on those, as shown in Listing 11

```
import scala.util.Random
(Random.nextString(Random.nextInt(16)), Random.nextInt(100)) match {
  Matches tuples with any first element, and second element == 10.
  case (s, 10)                    ⇒ The first element is accessible as s here.

  Matches only the ("Lucky?", 42) tuple.
```

```
  case ("Lucky?", 42)                        ⇒ None of the elements are accessible here.

  Matches any tuple as long as the first element is longer than 5 characters and second element is less than 10.
  case (x, y) if x.length > 5 && y < 10 ⇒ Both first and second elements are accessible as x and y.

  Matches any value
  case _                                     ⇒ None of the elements are accessible here.
}
```

Listing 11: Pattern matching II

As useful as tuples are for quick ad-hoc structures, the real deal in pattern matching are case classes. Just like matching on a tuple, it is possible to match on case classes and extract their fields as needed. In addition to using the ***match*** keyword followed by ***case*** s, pattern matching also applies to declarations of variables in Listing 12.

```
Pattern matches on the right-hand side, pulls out the first parameter as val first: String.
val Person(first, _, _) = Person("Foo", "Bar", 99)

Pattern matches on the right-hand side, comparing the value of the first parameter to be equal to in-scope variable first. (Notice the backticks.)
val Person(`first`, _, _) = Person("Foo", "Bar", 99)

Because there are no alternatives, if the val pattern match fails, it raises an exception.
val Person(`first`, _, _) = Person("Fooz", "Bar", 99)
```

Listing 12: Pattern matching III

The pattern match that raises an exception if it fails is often convenient to use in tests. Suppose there

Exercise: Using case classes and pattern matching

Build an numerical expression evaluator; the evaluator should support binary operators $+, -, \times, /$ and support arbitrary nesting. It should be able to evaluate, for example $4 + 3(5/(12 - 7))$. To make things simpler, it is not necessary to include error reporting (division by zero, for example), and it is not necessary to provide a way to turn a *String* into an expression.

```
The sum type Expr's subtypes define the "operations" we support.
sealed trait Expr
The different cases need to support our operations and constants.
case class Plus(left: Expr, right: Expr) extends Expr
...
case class Const(const: Int) extends Expr

It will be convenient to package our evaluator in its own module.
object Evaluator {

  Hint: you will need to pattern-match on the different Exprs.
  def eval(expr: Expr): Int = expr match {
    case Const(x) ⇒ x
    ...
  }

}

object EvaluatorMain extends App {

  val e = Evaluator.eval(Plus(Plus(Const(4), Const(5)), Const(3)))
  println(e)  | 12

}
```

Listing 13: Expression evaluator

Exercise (extra): logical expression simplifier

Build a logical expression simplifier that reduces the number of logical operations to be performed, and can spot tautologies and contraditions. For example, $(a \wedge b) \vee (a \wedge b)$ should simplify to just $a \wedge b$; $(a \vee \neg a)$ should simplify to $\top$ (tautology; always true); $(a \vee \neg a) \wedge (a \oplus a)$ should simplify to $\bot$ (contradiction; always false).

```
The sum type LExpr defines the different logical expressions we support.
sealed trait LExpr

The different cases need to support our operations and constants.
case class LAnd(left: LExpr, right: LExpr) extends LExpr
case object Contradiction extends LExpr
...
case class Var(named: String) extends LExpr

Just like the evaluator, package the simplifier in its own module.
object Simplifier {
```

```
    Hint: you will need to pattern-match on the different LExprs.
    def simplify(expr: LExpr): LExpr = expr match {
      ...
    }

  }

  object SimplifierMain extends App {

    val se = Simplifier.simplify(
      LAnd(
        LOr(Var("a"), LNot(Var("a"))),
        LXor(Var("a"), Var("a"))
      )
    )
    println(se)  | "Contradition"

  }
```

Listing 14: Logical expression simplifier

2.4 For comprehensions

TODO

2.5 Implicits

Implicits are one of the key features of Scala. As the word suggests, implicit parameter values are supplied "automatically" by the compiler looking for the appropriate values in the current *implicit scope*[3]. The look-up ignores the parameter names, it only cares about the types. Implicits also apply when the compiler encounters an identifier on a value of some type that the type does not implement. In that case, the compiler will follow the same implicit scoping rules to find an implicit conversion that turns the given type into another type that contains the identifier.

```
object SimpleImplicitsMain extends App {
  Defines an implicit conversion from a String into LoudString, which adds the method LOUD.
  implicit class LoudString(s: String) {
    def LOUD: String = s.toUpperCase() + "!!"
  }

  Define the implicit value of type String.; notice the implicit application of the LoudString.LOUD.
  implicit val completelyArbitraryName: String = "Hello, world".LOUD

  The sayHello1 method has single implicit parameter list, with one parameter of type String.
  def sayHello1(implicit greeting: String): Unit = println(greeting)

  The sayHello2 method has two parameter lists, one empty, and one implicit with one parameter of type String.
  def sayHello2()(implicit greeting: String): Unit = println(greeting)

  sayHello1      | to apply, omit the implicit parameter list, leaving just sayHello1.
  sayHello2()    | to apply, omit the implicit parameter list, leaving just sayHello2().
}
```

Listing 15: Simple implicits

Defining implicit `String` s demonstrates the principle, but it is rather useless because values of type `String` are so ubiquitous in typical programs. Different types are sometimes more useful; consider a `javax.sql.Connection` which might be passed implicitly to various data access code to avoid having to do too much typing. The ability to turn values into values that contain identifiers that are not available on the original types are also quite useful. The real power, though, comes from the fact that the implicit resolution does not stop at the first step. The compiler follows all possible paths, using as many conversions as necessary, to get the code to compile.

Pimp my library. Providing implicit conversions from existing types to other types that contain useful or convenient identifiers is called *pimp my library*. It is a fancy name for the adapter pattern[1], but implicits provide the convenience to avoid having to create instances of the adapters manually. It is a useful way to construct domain-specific languages, particularly when combined with Scala's ability to call single parameter method without dots and braces, as though it were an operator.

[3]The implicit scope is somewhat complex; for now, it will be sufficient to remember that implicit scoping is similar to regular visibility scoping

```scala
implicit class RichDouble(x: Double) {
  def ^(y: Double): Double = math.pow(x, y)
}

val x: Double = 42
x ^ 2  | Computes the second power of x by essentially doing new RichDouble(x).^(2).
```

Listing 16: Pimp my library

Exercise: Pimp my library

Build a DSL that allows the expression evaluator to be written using natural-looking Scala code–though sadly not using the standard $+, -, *, /$ operators, but using our own *plus*, *minus*, *mult*, *div* instead. The expression *5 plus 10* should evaluate to *Plus(Const(5), Const(10))*. The type *Int* and *Expr* do not contain the methods *plus*, ..., *div*; it will be necessary to implement implicit conversions (use ***implicit class*** for convenience) that contains those methods.

```scala
It will be useful to package the PML conversions in their own module.
object Expr {
  Two conversions will be necessary: Int → Expr, Expr → RichExpr. The two P.M.L. implicit classes
  will contain the same functions; the only exception is that the one for Int will need to box the given
  Int into Const.
  implicit class RichInt(x: Int) {
    def plus(that: Expr): Expr = Plus(Const(x), that)
    ...
  }
  implicit class RichExpr(x: Expr) {
    def plus(that: Expr): Expr = Plus(x, that)
    ...
  }

}

object EvaluatorMain extends App {
  import Expr._
  val e = Evaluator.eval((5 plus 10) minus 8)
  println(e)
}
```

Listing 17: Expression evaluator

Hints:

- because the methods *plus(r: Expr): Expr*, ..., *div(r: Expr): Expr* would be the same, consider defining them in a trait (***trait*** *ExprOps* would be a jolly good name!) that is then mixed into the implicit classes, where the trait defines abstract member ***self****: Expr*, which is then used as the left-hand side in the *Expr* data constructors; the right-hand side coming from the parameter of the method.
- implementing a class that adds the methods *plus*, ..., *div* to *Int* will allow for code such as *10 plus Const(5)*, but it will not compile *10 plus 5*, because there is no function *plus*, ..., *div* that takes an *Int*. The preferred option is to add those over implicit conversion from *Int* to *Expr*: ***implicit def*** *intToExpr(i: Int): Expr* is generally frowned-upon because it allows silent and potentially very powerful conversions.

If you are feeling inventive or perhaps mischievous, pick Unicode identifiers for the boring ASCII method names in the DSL. Resurrect APL by using ÷ instead of *div*!

There are other languages that offer similar functionality. Swift's extensions, particularly with type constraints, certainly feel like typeclasses.

Type classes. Type class is a concept from Haskell[3]; it is a definition of methods that can be implemented for some type. A type class is therefore an interface parametrized by a type; the implementations are instances of this interface for some types. Suppose the expression evaluator needs to be able to evaluate not just *Int* s, but also other number-like values. The only change to its structure is an interface that implements the addition, subtraction, multiplication, and division for *some* number-like type; see Listing 18.

```scala
Define the typeclass that defines the methods that instances for the type A must implement.
trait NumberLike[A] {
  def plus(x: A, y: A): A
  ...
  def div(x: A, y: A): A
}
```

```scala
sealed trait Expr
...
Modify the Const data constructor to take any type A instead of the concrete Int.
case class Const[A](a: A) extends Expr

object Evaluator {

  The eval function now needs to take the interface that implements the number-like behaviour for type A.
  def eval[A](expr: Expr, numberLike: NumberLike[A]): A = expr match {
    case Plus(l, r) ⇒ numberLike.plus(eval(l, numberLike), eval(r, numberLike))
    ...
    case Const(a: A) ⇒ a
  }

}

object EvaluatorMain extends App {
  import Expr._
  import Evaluator._

  val expr = 5 plus 10
  eval(expr, new NumberLike[Int] {
    def plus(x: Int, y: Int): Int = x + y
    ...
    def div(x: Int, y: Int): Int = x / y
  })  | evaluates to 15.

}
```

Listing 18: Expression evaluator without typeclasses

This is rather tedious to write, especially all the points where the *numberLike* instance has to be passed around to the recursive calls of *eval*. Moving it to the implicit parameter list makes the code cleaner in Listing 19.

```scala
trait NumberLike[A] { ... }

object Evaluator {

  def eval[A](expr: Expr)(implicit N: NumberLike[A]): A = expr match {
    case Plus(l, r) ⇒ N.plus(eval(l), eval(r))
    ...
    case Const(a: A) ⇒ a
  }

}

object EvaluatorMain extends App {
  import Expr._
  import Evaluator._

  val expr = 5 plus 10
  eval(expr)(new NumberLike[Int] {...})  | evaluates to 15.

}
```

Listing 19: Expression evaluator with typeclasses-ish

This is an improvement, though it is still annoying to have to explicitly specify the value of the *numberLike* parameter in the implicit (sic!) parameter list in *EvaluatorMain*. The Scala compiler is able to find the implicit value, if one is available in the implicit scope; all that remains to be done is to provide an *instance of the NumberLike typeclass for the type Int*, and any other required types. (Viz Listing 20.)

```scala
trait NumberLike[A] { ... }

object EvaluatorMain extends App {
  import Expr._
  import Evaluator._

  implicit object IntNumberLike extends NumberLike[Int] {
    def plus(x: Int, y: Int): Int = x + y
    ...
  }

  Complex numbers are numbers too!
  case object Complex {
    val e: Complex = ...
    val pi: Complex = ...
    val i: Complex = ...
  }
  case class Complex(re: Double, im: Double) {
    def +(rhs: Complex): Complex = ...
```

```
    def ^(rhs: Complex): Complex = ...
  }
  implicit object ComplexNumberLike extends NumberLike[Complex] {
    def plus(x: Complex, y: Complex): Complex = x + y
    ...
  }

  And now the evaluator works for any type A that lies in the NumberLike typeclass; i.e. where there is an in-scope
  way to access or follow steps to create instances of NumberLike[A].
  eval[Int](5 plus 10)          | evaluates to 15.
  eval[Complex](e ^ (pi * i))   | evaluates to -1.

}
```

Listing 20: Expression evaluator with typeclasses

Exercise: Flexible evaluator

Allow the expression evaluator to operate on any number-like types, not just *Int* s. Instead of relying on inheritance, forcing all users of the evaluator to evaluate values conforming to some trait (with the number-crunching methods *+, −, ∗, /*), use parametric polymorphism and type classes. The *Const* data constructor will need to take any type instead of *Int* ; *RichInt* implicit class will need to pimp any type, not just *Int* ; the *eval* method will also need to be generic, but will need to (implicitly) require a typeclass: use the *Fractional* provided by the standard library, or define your own *NumberLike* typeclass. (The Scala standard library defines *Fractional* and *Integral* type classes, which contain the division operation whereas *Numeric* typeclass does not. Pick one, the limitation that not all types are fractional or integral. “Dotty will fix this!”)

```
object Expr {

  trait ExprOps {
    def self: Expr
    def plus ...  This will need to take generic type A.
    ...
  }

  implicit class RichA ...  This will need to take generic type A.
  implicit class RichExpr(val self: Expr) extends ExprOps

}

sealed trait Expr
...
case class Const ...  This will need to take generic type A.

object Evaluator {

  It will be convenient to use the full implicit syntax here
  def eval[A](expr: Expr)...: A = expr match {
    ...
  }

}

object EvaluatorMain extends App {
  import Expr._
  The point of application constrains the generic parameter, it is necessary to specify the type A. in eval.
  val e = Evaluator.eval[Double]((5.4 plus math.Pi) minus 8.8)
  ...
}
```

Listing 21: Flexible evaluator with DSL and typeclasses

Hints:

- it is not necessary to constrain the type of the *Const* data constructor; it can be “forall A.” *Const[A]* .
- it will be convenient to use the full implicit syntax in the *eval* method: ***def*** *eval[A](expr: Expr)(* ***implicit*** *N: Fractional[A]): A* instead of ***def*** *eval[A : Fractional](expr: Expr): A*
- the type information for the generic type *A* will be erased, causing compiler warning in the pattern match. It is possible to eliminate the warning by using another type class. Look for *Scala type tags and manifests.*

Unlike subtype polymorphism, typeclasses bring *parametric polymorphism*. Parametric polymorphism is a different way of requiring values to contain specified behaviour. Subtype polymorphism for the evaluator would require the values in the *Const* data constructor to be some supertype of all numbers. That does not sound so bad until one realises that the methods *plus*, ..., *div* in this supertype would have to return that supertype. Imagine in Listing 22 that it is possible to extend the types in *scala.lang*, like *Int*.

```
trait Num {
  def plus(that: Num): Num
  def minus(that: Num): Num
  ...
}

This isn't actually allowed, but humour the author.
class MyInt extends Int with Num {
  Attempts to implement the Num trait won't be successful...
  def plus(that: Num): Num = this + that
                                    ↑ this doesn't quite work: can't add Int and Num.

  Or the implementation won't at all be satisfying.
  def plus(that: Num): Num = that match {
    case x: MyInt ⇒ this + x
    case _ ⇒
        ↑ now what? throw exception?; add toInt() method to Num?
  }
}

This is Scala's common approach to static of-like methods in Java. Recall Optional.of and similar.
object MyInt {
  def apply(i: Int): MyInt = new MyInt(i)
}

But never mind all above, the show must go on.
case class Const(a: Num) extends Expr

def eval(expr: Expr): Num = expr match {
  case Plus(l, r) ⇒ eval(l).plus(eval(r))
  ...
}

I allowed Int to be extended; requiring integer literals to be instances of MyInt is too much!
Using the object MyInt.apply allows us to avoid the new keyword in this case.
val e = Evaluator.eval(MyInt(5) plus MyInt(3) minus MyInt(1))
     ↑ the type of e is Num, even though MyInt was used.
```

Listing 22: Exploring the Num supertype

The code in Listing 22 is bad enough *now*, but it has the potential to become much *worse* if it becomes widely used. The authors of the code were able to implement sensible inheritance structure for *MyInt*s, *MyDouble*s, even *MyBigDecimal*s and shipped the JAR. Then someone else decided to use the expression evaluator for complex numbers; but even with the implementation of ***class*** *Complex* ***extends*** *Num {...}*, it is still possible to call *MyInt(5).plus(Complex(1, 4))*. The class *MyInt* has no knowledge about the *Complex* class! (Exactly what happens in that call depends on the implementation of the *MyInt.plus* method, but it won't be pretty.)

Parametric polymorphism, as demonstrated in your own work on the Flexible evaluator exercise, presents a neat solution. It is possible to use typeclasses to add behaviour for concrete types without losing the concreteness.

A remark on Haskell. Scala's take on typeclasses should clarify Haskellers' usual statement that "typeclasses are not like OOP classes, they are more like interfaces." Quite: a typeclass is an interface with a type parameter that defines behaviour for the type. Its implementations specify a concrete type and implement the behaviour. Typeclass instances are then regular values (remember the singletons from the Scala code?) that implement the typeclass interface.

```
Require the type A to have an instance of the NumberLike trait available for it.
Compiler will supply the value in the parameter N of the implicit parameter list.
def eval[A](expr: Expr)(implicit N: NumberLike[A]): A = expr match {
  Call the typeclass's methods using the standard method call notation.
  case Plus(l, r) ⇒ N.plus(l, r)
  ...
}

Define the typeclass NumberLike for type a with methods like plus, ... .
class NumberLike a where
  plus :: a → a → a
  ...

Provide an instance of NumberLike for the type Int.
instance NumberLike Int where
  plus x y = ‘addl %rdx, %rax‘  | Hoping that we have fastcall convention and the parameters are in %rax and %rdx.
  ...

Provide an instance of NumberLike for the type Complex.
instance NumberLike Complex where
```

```
...
```

Just like the Scala version, `eval` requires the type `a` to lie in the `NumberLike` typeclass; that is, that there is an instance of `NumberLike` for that `a`.

```
eval :: (NumberLike a) ⇒ Expr → a
```

However, the typeclass instance does not get a name in the function. Instead, all its functions `plus`, ... are available in the function without any additional notation.

```
eval (Plus l r) = plus (eval l) (eval r)
                  ↑
                  In Scala, this would be N.plus(eval(l), eval(r))
```

2.6 Rich type system.

Variance. Kinds & Type constructors. Higher-kinded types.

Exercise: Flexible evaluator II

Explore and explain the difference in just adding the generic parameter to the *Const* data constructor (leaving ***sealed trait*** *Expr* intact), or letting *Expr* become ***sealed trait*** *Expr[A]*. Explain the effects of covariant type parameter in ***sealed trait*** *Expr[+A]*. (Try ***val*** *e: BigDecimal = Evaluator.eval(Const(4))* with invariant *Expr* and then with covariant *Expr* to get started.)

Listing 23: Variance

3 The standard library

The Scala standard library defines many standard structures in case classes. Picking the two most useful ones, *Option* and *Either* are defined as case classes. This means that it is possible to pattern match on their values.

These are the definition in the standard library; note that the actual definitions are more complex, but this code expresses the essence.

`Option[A]` has two alternatives: `Some[A]` and `None`. It is used to express one-or-missing values.

```
sealed trait Option[A]
case class Some[A](a: A) extends Option[A]
case object None extends Option[Nothing]
```

`Either[L, R]` has also two alternatives: `Left[L]` and `Right[R]`. It is used to express success or failure, with success by convention on the right and failure on the left.

```
sealed trait Either[L, R]
case class Right[R](r: R) extends Either[Nothing, R]
case class Left[L](l: L) extends Either[L, Nothing]
```

Listing 24: Pattern matching in stdlib

...

4 Spring Framework

The Spring Framework is a dependency injection framework; it encourages composition over inheritance, it encourages expressing dependencies as interfaces rather than concrete implementations. The framework takes care of instantiating the components in the correct order; most components (the ones that fall into the *@Component* stereotype) are *singletons*. This means that it is possible to treat the *@Component*-annotated components as namespaces rather than containers of state[4]. The reason why Spring Framework encourages programming to interfaces is to make the software easily testable: there can be separate implementations or mocks for unit and integration tests.

```
interface ReportGenerator {
    Generates the PDF report for the given user,
    returns the byte array representing the PDF contents
    byte[] generate(final String user);
}

@Component
public class ReportService {
    private final ReportGenerator reportGenerator;
```

[4] In fact if the methods in *@Component*-annotated classes mutates & accesses its fields, it will suffer from race conditions.

```
    @Inject
    public ReportService(final ReportGenerator reportGenerator) {
        this.reportGenerator = reportGenerator;
    }

    public void reportAll() {
        for (final String user : Arrays.asList("a", "b", "c")) {
            final byte[] pdf = this.reportGenerator.generate(user);
            Now you're on your own...
        }
    }
}
```

Listing 25: Components

For a Spring Framework application to be able to construct the *ReportService*, it needs to be able to construct exactly one component that implements the *ReportGenerator* interface.

```
@Component
public class JasperReportsReportGenerator implements ReportGenerator {
    public byte[] generate(final String user) {
        ...
    }
}
```

Listing 26: Components

Without a DI framework, the work of constructing the dependencies would fall on the programmers, yielding code similar to Listing 27.

```
Typically in public static void main(String[] args) or in a test:

ReportGenerator rg = new JasperReportsReportGenerator();
ReportService rs = new ReportService(rg);
```

Listing 27: Manual DI

Constructing the instances of the *JasperReportsReportGenerator* and *ReportService* using their constructors isn't a problem per se, but with growing number of dependencies this grows to be tedious.

5 Zero to hundred

FizzBuzz is a typical program that follows *Hello, world*, adding iteration and conditions. The Scala version of FizzBuzz is shown in Listing 28–it shows the definition of a function ***def***, followed by name and arguments, and its implementation that follows the = sign. The loop (***for***) and condition (***if*** , ***else***) keywords are the old friends from other languages.

```
def fizzBuzz = {
  for (i ← 1 to 100) {
    if (i % 15 == 0) println("FizzBuzz")
    else if (i % 3 == 0) println("Fizz")
    else if (i % 5 == 0) println("Buzz")
    else println(i)
  }
}
```

Listing 28: Fizz Buzz

The FizzBuzz from Listing 28 isn't particularly re-usable: it simply prints 100 elements to the standard output, nothing else and nothing more. There is no way, for example, to direct the output to a web socket, or to use it to determine how it maps of the value in the *integer* domain to the "FizzBuzz domain". Hmm!–*mapping* and *domain* sound like mathematics; and functional programming is supposed to be somehow more mathematical. And mathematics is jolly wonderful.

The first step in making the *fizzBuzz* more mathematical is to make it map an input to exactly one useful output. Right now, its return type now is *Unit*, which is a bit like *void* in Java and C; changing its definition to ***def*** *fizzBuzz2(max: Int): Unit* (and then using the *max* parameter in the loop) isn't particularly useful: it is a mapping from a number to *Unit*. And, if this were mathematics, there can be only one such mapping: ***def*** *fizzBuzz2(max: Int): Unit = ()*. Instead of printing the elements to the console, the implementation needs to return a value that can be printed. A a simple *String* would do, but a *Seq* of *String* s is better. The type becomes *Int ⇒ Seq[String]*, and the implementation is shown in Listing 29.

```
def fizzBuzz(max: Int): Seq[String] = {
  var result = List.empty[String]
  for (i ← 1 to max) {
```

```
    if (i % 15 == 0) result = result :+ "FizzBuzz"
    else if (i % 3 == 0) result = result :+ "Fizz"
    else if (i % 5 == 0) result = result :+ "Buzz"
    else result = result :+ i.toString
  }
  result
}
```

Listing 29: Fizz Buzz

This is a huge improvement! The *fizzBuzz* is now indeed a function: it maps input to output and its result depends only on the value of the parameter. It would even be possible to pre-compute the result for all possible values of the input and replace the function's body with a look-up in that table: the function would become just data!

Well, the outside looks great, but the implementation stinks! It uses mutation, and what about the strange `:+` operator in `result :+ "Fizz"`, never mind the `for (i ← 1 to max) {...}` nonsense!

```
def fizzBuzz(max: Int): Seq[String] = {
  def fb(i: Int): String =
    if (i % 15 == 0) "FizzBuzz"
    else if (i % 3 == 0) "Fizz"
    else if (i % 5 == 0) "Buzz"
    else i.toString

  (1 to max).map(fb)
}
```

Listing 30: Fizz Buzz

In Scala, every concrete type (except `Nothing`) can have a value: for example, the type `Boolean` is inhabited by values `true`, `false`; the type `Int` is inhabited by values such as `5, 42, -100, 0, ...`; the type `String` is inhabited by values such as `"Hi", ":)", ""`; the type `Unit` is inhabited by the only value `()`. (No, really, it's perfectly good Scala syntax to write `()` as value. It's just not particularly useful.) The only type that does not have any inhabitants is `Nothing`: it represents expressions that *diverge*, for example throwing an exception.

Taking a more precise look at `def fizzBuzz` reveals its type to be `Unit`; it evaluates to only one value, namely `()`. If it were a function in the sense of strictly mapping input to output, it would be no different from any other `()` *constant*. But `fizzBuzz` does some additional work before returning `()`; this additional work is not represented by its type, even though it is its *raison d'être*.

In Java and C, there is no *value* of type `void`

As it stands, its type is `() ⇒ Unit`,

6 Pattern matching

sasd

7

References

[1] *Adapter Pattern*. https://en.wikipedia.org/wiki/Adapter_pattern.

[2] *F#*. https://fsharp.org.

[3] *Haskell*. https://www.haskell.org.